

WSE1 – Werkzeuge im Software Engineering

SE.ba VZ

Docker Crash Course

Johannes Karder, Andreas Scheibenpflug, Sebastian Pimminger

The Birth of Docker



<https://imgur.com/3eTKEZp>

Docker

- Set of tools to ...
 - ▶ create containers and images
 - ▶ configure them (Dockerfile)
 - ▶ execute them (Docker engine)
- Docker Hub
 - ▶ Platform for image distribution



The screenshot shows the Docker Hub search interface. At the top, there's a search bar with the placeholder "Search Docker Hub". To the right are links for "Explore", "Pricing", "Sign In", and a "Sign up" button. Below the search bar, the text "1 - 25 of 10.000 available results." is displayed, along with a dropdown menu set to "Suggested". On the left, there are "Filters" for "Products" (Images, Extensions, Plugins), "Trusted Content" (Docker Official Image, Verified Publisher, Sponsored OSS), and "Operating Systems" (Linux). Two search results are shown: "alpine" and "nginx".

- alpine** Docker Official Image · 1B+ · 10K+
Updated 5 days ago
A minimal Docker image based on Alpine Linux with a complete pack...
Linux IBM Z riscv64 x86-64 ARM 386 PowerPC 64 LE ARM 64
- nginx** Docker Official Image · 1B+ · 10K+
Updated 4 days ago
Official build of Nginx.
Linux 386 mips64le PowerPC 64 LE IBM Z x86-64 ARM ARM 64

<https://hub.docker.com/search?q=%E2%80%A2>

Container 1/2



container noun

con·tain·er kən-'tā-nər

: one that **contains**: such as

- a** : a receptacle (such as a box or jar) for holding goods
- b** : a portable compartment in which freight is placed (as on a train or ship) for convenience of movement


<https://www.merriam-webster.com/dictionary/container>

Container 2/2


- A container ...
 - ▶ is an isolated environment
 - ▶ is managed by a host operating system (OS)
 - ▶ executes processes
- These processes ...
 - ▶ have dedicated resources (file system, files, devices, libraries)
 - ▶ cannot directly access resources of the host OS

Virtual Machines vs. Containers

Virtual machines




Containers




<https://www.atlassian.com/microservices/cloud-computing/containers-vs-vms>

Overview



<https://github.com/docker/docs/blob/6a66f748791143ecfe4c321f9e11fd4b641f8948/content/engine/images/engine-components-flow.png>

Docker Architecture



<https://github.com/docker/docs/blob/6a66f748791143ecfe4c321f9e11fd4b641f8948/content/engine/images/architecture.svg>

Docker Client

- Frontend for users
 - ▶ Manage containers, volumes, networks, ...
 - ▶ Create images using Dockerfiles
 - ▶ Download images from registry
 - ▶ ...
 - ▶ `docker --help`
- Uses REST API to communicate with Docker daemon
- `docker run debian /bin/echo "Hello, Docker world!"`

Docker Daemon

- Creates and manages containers and images
- Provides REST API
 - ▶ Programmatically use Docker
 - ▶ Used by Docker CLI
- Container runtime: containerd + runc
- Alternative container runtimes:




crun

<https://github.com/containers/crun>



youki

<https://github.com/containers/youki>



Docker Images and Containers


- Preconfigured runtime environment
- Technically, an archive that contains ...
 - ▶ one or more file systems with files and
 - ▶ metadata
- `docker run` starts containers from images
- A container is a runnable instance of an image
- Images should be considered immutable

Docker Registry

- Service to provide Docker images
- Official Docker registry: <https://hub.docker.com/>
- Self-hosted registry: https://hub.docker.com/_/registry
- Image references
 - ▶ NAME[:TAG | @DIGEST]
 - ▶ name → debian:latest
 - ▶ name:tag → debian:bullseye
 - ▶ name@digest → debian@sha256:e538a2f...

Dockerfiles


- Contain instructions for creating images
- `docker build` creates an image from a Dockerfile



Copy-on-write Storage


- Union filesystem: Allows to mount and overlay multiple file systems
- Docker: Overlay multiple layers of a docker image
- Read operations executed on existing layers
- Write operations executed on new layer
 - ▶ File to be written to is copied first → Copy-on-write
- Docker provides multiple storage drivers
 - ▶ OverlayFS is default on currently supported Linux distributions

Image Layers



<https://docs.docker.com/storage/storagedriver/>

Container Layers



<https://docs.docker.com/storage/storagedriver/>

Shared Layers

debian:latest

```
FROM scratch
ADD rootfs.tar.xz /
CMD ["bash"]
```

custom:latest

```
FROM debian:latest
LABEL author='johannes.karder@fh-ooe.at'
COPY ./test.txt /root/test.txt
RUN echo 'hello' > /root/hello.txt
CMD ["date"]
EXPOSE 80
```

Layers (2)

| | | |
|---|---|-----------|
| 0 | ADD file:b4987bca8c4c4c640d6b71dccfd7172b44771e0f851a47d05c00c2bdcd204f6 in / | 116.51 MB |
| 1 | CMD ["bash"] | 0 B |

Layers (7)

| | | |
|---|---|-----------|
| 0 | ADD file:b4987bca8c4c4c640d6b71dccfd7172b44771e0f851a47d05c00c2bdcd204f6 in / | 116.51 MB |
| 1 | CMD ["bash"] | 0 B |
| 2 | LABEL author=johannes.karder@fh-ooe.at | 0 B |
| 3 | COPY ./test.txt /root/test.txt # buildkit | 16 B |
| 4 | RUN /bin/sh -c echo 'hello' > /root/hello.txt # buildkit | 6 B |
| 5 | CMD ["date"] | 0 B |
| 6 | EXPOSE map[80/tcp:{}] | 0 B |

docker build and Caching

- docker build creates a new layer for various commands within a Dockerfile
- For file operations (e.g. COPY) the checksum of the files is saved
- All layers are cached
 - ▶ Speeds up build process
 - ▶ Allows sharing of layers between different images
 - ▶ Speeds up image downloads
 - ▶ Reduces size of containers and images

Docker Commands

Build Images

- `docker build <path>`
- Builds an image from a Dockerfile and saves it locally
- Options
 - ▶ `-f / --file`: Specify Dockerfile (default is “PATH/Dockerfile”)
 - ▶ `-t / --tag`: Name and optionally tag an image
 - ▶ `docker build -t myimage:latest .`

List Images

- docker images
- docker image ls
- Lists all local Docker images, e.g. saved in /var/lib/docker/image/
- Options
 - ▶ -a / --all: Also list intermediate images

Save / Load Images

- `docker [image] save <image> > image_file.tar`
- `docker [image] save <image> -o image_file.tar`
 - ▶ Saves image as tar archive

- `docker [image] load < image_file.tar`
- `docker [image] load -i image_file.tar`
 - ▶ Loads image from tar archive

Remove Images

- `docker rmi <image>`
- `docker image rm <image>`
 - ▶ Deletes an image

Pull / Push Images

- docker [image] pull <[url/]tag>
 - ▶ Loads an image from a repository
- docker [image] push <[url/]tag>
 - ▶ Uploads an image to a repository

Run Containers 1/2

- `docker run <name[:tag]> [command] [args]`
- Starts an image, and optionally a command within the resulting container
- Creates a writable layer
- Options
 - ▶ `-i / --interactive`: Leaves open STDIN even if not attached
 - ▶ `-a / --attach`: Attach to STDIN/STDOUT/STDERR
 - ▶ `-t / --tty`: Opens a shell for input/output
 - ▶ `-v / --volume`: Mounts directories/files/volumes

Run Containers 2/2

- Options
 - `-p` / `--publish`: Opens ports of the container to the host OS
 - `-d` / `--detach`: Starts container in the background
 - `--rm`: Automatically removes container when it exits
 - `-e` / `--env`: Sets environmental variables within the container
 - `--name`: Sets a name (can be used in other commands, e.g. start/stop/etc.)
 - Many more options to constrain resources
 - `-c` / `--cpu-shares`
 - `--cpus`
 - `-m` / `--memory`

Create Containers

- `docker container create [options] <image> [command] [args]`
- Creates a writable layer for a container on top of an image
- Can be started with `docker start`
- Options: see `docker run`

Start / Stop


- `docker stop <container>`
 - ▶ Stops a running container
 - ▶ Sends SIGTERM to all container processes and SIGKILL after timeout
- `docker kill <container>`
 - ▶ Sends SIGKILL to all container processes
- `docker start <container>`
 - ▶ Starts a stopped container
- `docker pause <container>`
 - ▶ Pauses processes within the container (cgroups freezer)
- `docker unpause <container>`

Copy To / From Containers

- Copy directories/files between container and host
 - ▶ `docker [container] cp <container>:src dest`
 - ▶ `docker [container] cp src <container>:dest`
- Example
 - ▶ `docker cp container1:/root/notes.txt notes.txt`

Mounts

- Allow to
 - ▶ ... persist data (volumes)
 - ▶ ... access/exchange data with host system (bind mounts)
 - ▶ ... temporary store data (tmpfs)



<https://docs.docker.com/storage/volumes/>

Volumes

- Managed by Docker
- Mounted in containers
- Data remains persisted across container restarts
- Format: `volume_name:container_dir[:options]`
- Example:
 - ▶ `docker volume create my_volume`
 - ▶ `docker run -v my_volume:/root`

Volume

- docker volume
 - ▶ `create <name>`: Creates a named volume
 - ▶ `rm <name>`: Deletes a volume
 - ▶ `ls`: Lists created volumes
 - ▶ `prune`: Deletes all unused volumes

Bind Mounts

- Mount files or folders of host system in container
- Format:
 - ▶ `host_dir:container_dir[:options]`
 - ▶ `host_file:container_file[:options]`
- Example:
 - ▶ `docker run -v $PWD/notes.txt:/root/note.txt ...`
 - ▶ `docker run --mount \type=bind,\\source=$PWD/notes.txt,\\target=/root/note.txt ...`

Tmpfs

- Allows mounting in-memory file system
 - ▶ which is outside the writable layer of a container
- Tmpfs data is lost upon stopping the container
- Tmpfs mounts
 - ▶ ... cannot be shared between containers
 - ▶ ... only work on Linux
- Example:
 - ▶ `docker run -tmpfs /logs ...`

Ports

- Open container ports to host system
- Access applications bound to ports within container
- Format:
 - ▶ `host_ports:container_ports`
- Example:
 - ▶ `docker run -p 1234:80 ...`
 - ▶ `docker run -p 1234-1236:1234-1236 ...`

Ps

- docker ps lists all running containers
- Options:
 - ▶ -a / --all: list all containers, including the ones not running
 - ▶ -l / --latest: show last started container
 - ▶ -s / --size: show size of containers
 - Size: Size of writable layer
 - Virtual: Size of read-only image + writable layer

Attach / Exec

- `docker attach <container>`
 - ▶ Connects to STDIN, STDOUT and STDERR of container
- `docker exec <container> command [args]`
 - ▶ Executes command in running container

Top / Stats

- `docker top <container>`
 - ▶ Lists running processes of container
- `docker stats [<container>]`
 - ▶ Lists resource usage statistics of container(s)
 - ▶ CPU, memory, network I/O, disk I/O

Networking

- Multiple networking options
 - ▶ Extendable via plugins
- Default: bridge
 - ▶ Default network for containers – if not specified otherwise
 - ▶ No open ports
 - Between containers
 - Between container and host system
 - ▶ No name resolution between containers
 - Access only possible via IP addresses

User-Defined Bridges

- Assignment of containers to network
- Ports between containers are open
- Ports to host system must be opened explicitly
- DNS available
 - ▶ Container name == DNS name
- Container can be added/removed to/from user-defined bridges while running

Network

- `docker network create <network>`
 - ▶ Creates a new network (default: bridge)
- `docker network [dis]connect <network> <container>`
 - ▶ Adds/removes container to/from network
- `docker network rm|ls|prune`
 - ▶ See docker volume
- `docker run --net[work] <network> ...`
 - ▶ Starts a container within network
 - ▶ `docker run --network none ...` disables networking


Dockerfiles

Build Images

- `docker build <path>`
- Builds an image from a Dockerfile and saves it locally
- Options
 - ▶ `-f / --file`: Specify Dockerfile (default is “PATH/Dockerfile”)
 - ▶ `-t / --tag`: Name and optionally tag an image
 - ▶ `docker build -t myimage:latest .`

Dockerfiles

- Contain instructions for creating images
- `docker build` creates an image from a Dockerfile



FROM

- Defines the base image
 - ▶ FROM <image[:tag]> [AS <name>]
- Example:
 - ▶ FROM debian:latest
- Name is required for multi-stage builds, e.g.
 - ▶ COPY --from=<name> ...
- Multi-stage builds (> v17.05)
 - ▶ Multiple FROMs possible
 - ▶ Commands after last FROM make up final image

RUN

- Executes a command within image
- Used to configure image
 - ▶ Install software
 - ▶ Configure software
 - ▶ ...
- Examples:
 - ▶ RUN apt-get install nginx
 - ▶ RUN dotnet build

WORKDIR

- WORKDIR <path>
- Defines working directory for commands
 - ▶ Used to expand relative paths in COPY, RUN, CMD, ADD, ENTRYPOINT
- Example:
 - ▶ WORKDIR /root
 - ▶ COPY /root/file.txt .

COPY

- COPY <src> <dest>
- Copies files/directories from host system to image
 - ▶ Makes files available within container
 - ▶ Directories are created if non-existent
 - ▶ Multiple <src> may be defined
 - <dest> must then be a directory and end with a slash “/”
- Examples:
 - ▶ COPY ./notes.txt /root/
 - ▶ COPY ./NotesAPI/ .
 - ▶ COPY ./NotesAPI/ /root/NotesAPI/
 - ▶ COPY files* /root/
 - ▶ COPY ["a.txt", "b.txt", "/root/"]

ADD

- Like COPY, but
 - ▶ ... if <src> is an URL, downloads the file and saves it to <dest>
 - ▶ ... if <src> is a tar archive, extracts it in <dest>
- Examples:
 - ▶ ADD <https://planet.osm.org/planet/planet-latest.osm.bz2> /root/
 - ▶ ADD rootfs.tar.xz /

ENTRYPOINT

- Defines the executable to start once the container is started
- Formats:
 - ▶ Shell form: ENTRYPOINT command param1 param2 ...
 - ▶ Exec form: ENTRYPOINT ["executable", "param1", "param2", ...]
- Arguments to docker run <image> ...
 - ▶ ... are appended after all arguments in exec form
 - ▶ ... override all elements specified using CMD
- Override ENTRYPOINT using docker run --entrypoint ...

CMD

- Defines the program to be executed
- Defines default parameters for ENTRYPOINT
- Format:
 - ▶ Exec form: CMD ["executable", "param1", "param2", ...]
 - ▶ Shell form: CMD command param1 param2 ...
 - ▶ Parameter for ENTRYPOINT: CMD ["param1", ...]
- Exec form does not start a shell
 - ▶ No substitution, e.g. no shell variables are evaluated

ENTRYPOINT, docker run and CMD

- ENTRYPOINT in exec form:
 - ▶ Parameters of CMD are appended to ENTRYPOINT
 - ▶ Parameters of docker run are appended to ENTRYPOINT
- ENTRYPOINT in shell form:
 - ▶ Parameters of CMD / docker run are not appended
 - ▶ ENTRYPOINT is started as sub-process of /bin/sh
 - Application will not receive SIGTERM from docker stop

ENTRYPOINT - Example

```
FROM debian:stable
ENTRYPOINT ["/bin/echo"]
CMD ["hello, cmd"]
```

```
$ docker run ...           # prints "hello, cmd"
$ docker run ... "hello, docker" # prints "hello, docker"
```

ENV

- Sets environmental variables
- ENV <key>=<value>
- ENV <key> <value>
- Environmental variables are available in running container
- docker run --env <key>=<value>
- Examples:
 - ▶ ENV DOTNET_SDK_VERSION=1.0.4
 - ▶ ENV DOTNET_SDK_VERSION 1.0.4

ARG

- Defines variables that can be used during build time
- ARG <name>[=<default value>]
- Variables are not available in running container
- For multi-stage builds, they must be defined in every stage for them to be used
- User can specify values:
 - ▶ docker build --build-arg <key>=<value>
- Environmental variables override build variables

EXPOSE

- EXPOSE <port> [<port>/<protocol> ...]
- Only informs that container will listen on specified ports at runtime
- Does not publish port
- Must be done manually via docker run -p ...

Docker Compose

Docker Compose

- Definition, creation and execution of multi-container applications
- Definition via YAML file
- Specify required resources via configuration vs. docker run arguments
 - ▶ Port mappings
 - ▶ Networks
 - ▶ Volumes
 - ▶ ...

Structure

```
services:  
  service1:  
    ...  
  service2:  
    ...  
  
volumes:  
  ...  
  
networks:  
  ...
```

- One service corresponds to one Dockerfile/Image
- Configuration of each service entry corresponds to docker run/create/start arguments
- Sections for volumes and networks

docker compose

- docker compose build
 - ▶ Builds all images (executes build steps in YAML)
- docker compose up
 - ▶ Creates and starts all containers (services)
 - ▶ -d: Executes containers in background
- docker compose down
 - ▶ Stops all containers

`docker compose up`

- Creates containers at first start (`docker run`)
- Starts containers if they exist and are unchanged (`docker start`)
 - Writable layer remains as long as container exists
 - Can lead to hard to diagnose problems in docker compose files with many services
 - `--force-recreate`: always recreates containers

Example

```
services:  
  db:  
    container_name: mydb  
    image: postgres  
    volumes:  
      - ./tmp/db:/var/lib/postgresql/data  
  web:  
    build: .  
    command: bundle exec rails s -p 3000 -b '0.0.0.0'  
    volumes:  
      - .:/myapp  
    ports:  
      - "3000:3000"  
    depends_on:  
      - db
```

Further Reading

- Docker documentation
 - ▶ <https://docs.docker.com/get-started/docker-overview/>
- Dockerfile reference
 - ▶ <https://docs.docker.com/reference/dockerfile/>
- Compose file reference
 - ▶ <https://docs.docker.com/reference/compose-file/>